

Recommender Systems (I)

©Dr. Min Chi
mchi@ncsu.edu

The materials on this course website are only for use of students enrolled AIA and must not be retained or disseminated to others or Internet.



AI Academy

Examples:

amazon.com.




m o v i e l e n s
helping you find the *right* movies

last.fm™
the social music revolution

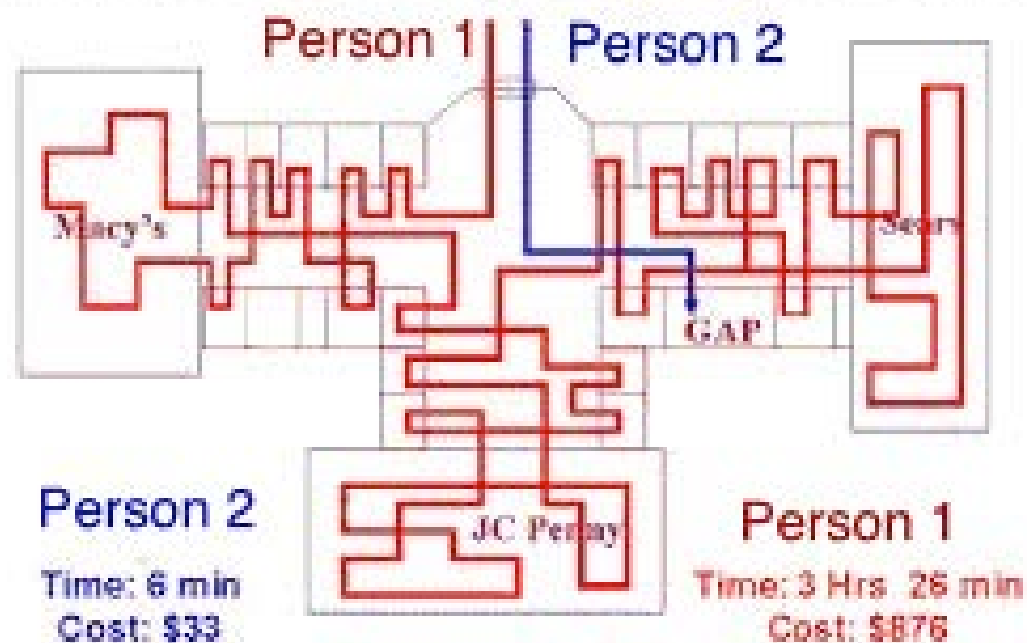
Google™
News

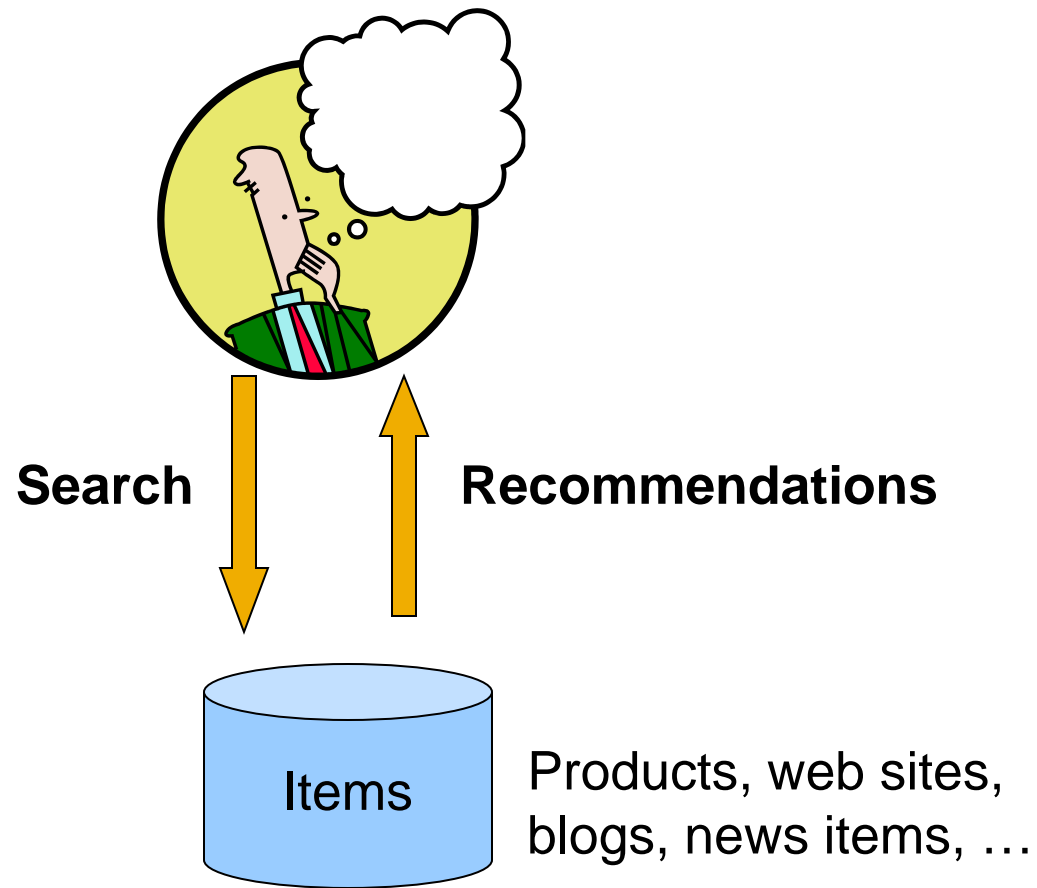
You Tube

XBOX
LIVE



Mission: Go to Gap, Buy a Pair of Pants





Examples:

amazon.com.



StumbleUpon



del.icio.us



m o v i e l e n s
helping you find the *right* movies

last.fm™
the social music revolution


Google™
News

You Tube

XBOX
LIVE

From Scarcity to Abundance

- **Shelf space is a scarce commodity for traditional retailers**
 - Also: TV networks, movie theaters,...
- **Web enables near-zero-cost dissemination of information about products**
 - From scarcity to abundance
 - Long Tail phenomenon.



Example

- Books, movies, music, news articles
- People (friend recommendations on facebook, LinkedIn, and Twitter)

More choice necessitates better filters

- Recommendation engines
- How "**Into Thin Air**" made **Touching the Void** a bestseller: <http://www.wired.com/wired/archive/12.10/tail.html>

Types of Recommendations

- **Editorial and hand curated**
 - List of favorites
 - Lists of “essential” items
- **Simple aggregates**
 - Top 10, Most Popular, Recent Uploads
- **Personalized to individual users**
 - Amazon, Netflix, ...

Terms

- X = set of **Customers**
- S = set of **Items**
- **Utility function** $u: X \times S \rightarrow R$
 - R = set of ratings
 - Ordered set
 - e.g., **0-5** stars, real number in $[0,1]$

Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Key Issues

- **Gathering “known” ratings for matrix**
 - How to collect the data in the utility matrix
- **Derive unknown ratings from the known ones**
 - Mainly interested in **high unknown** ratings
 - Not interested in knowing what you don't like but what you like
- **Evaluating methods**
 - How to measure success/performance of recommendation methods

1. Gathering Ratings

■ Explicit

- Ask people to rate items
- Doesn't scale: only a small fraction of users leave ratings and reviews.
- Crowdsourcing: Pay people to label items

■ Implicit

- Learn ratings from user actions
 - E.g., purchase implies high rating
- What about low ratings?

2. Deriving Unknown from Known

Key problem: Utility matrix U is **sparse**

- Most people have not rated most items
- **Cold start:**
 - New items have no ratings
 - New users have no history

Three approaches to recommender systems:

- 1) Content-based
- 2) Collaborative Filtering
- 3) Latent factor based

Content-based Recommendations

Main idea: Recommend items to customer x
similar to previous items rated highly by x

Example:

Movie recommendations

- same actor(s), director, genre, ...

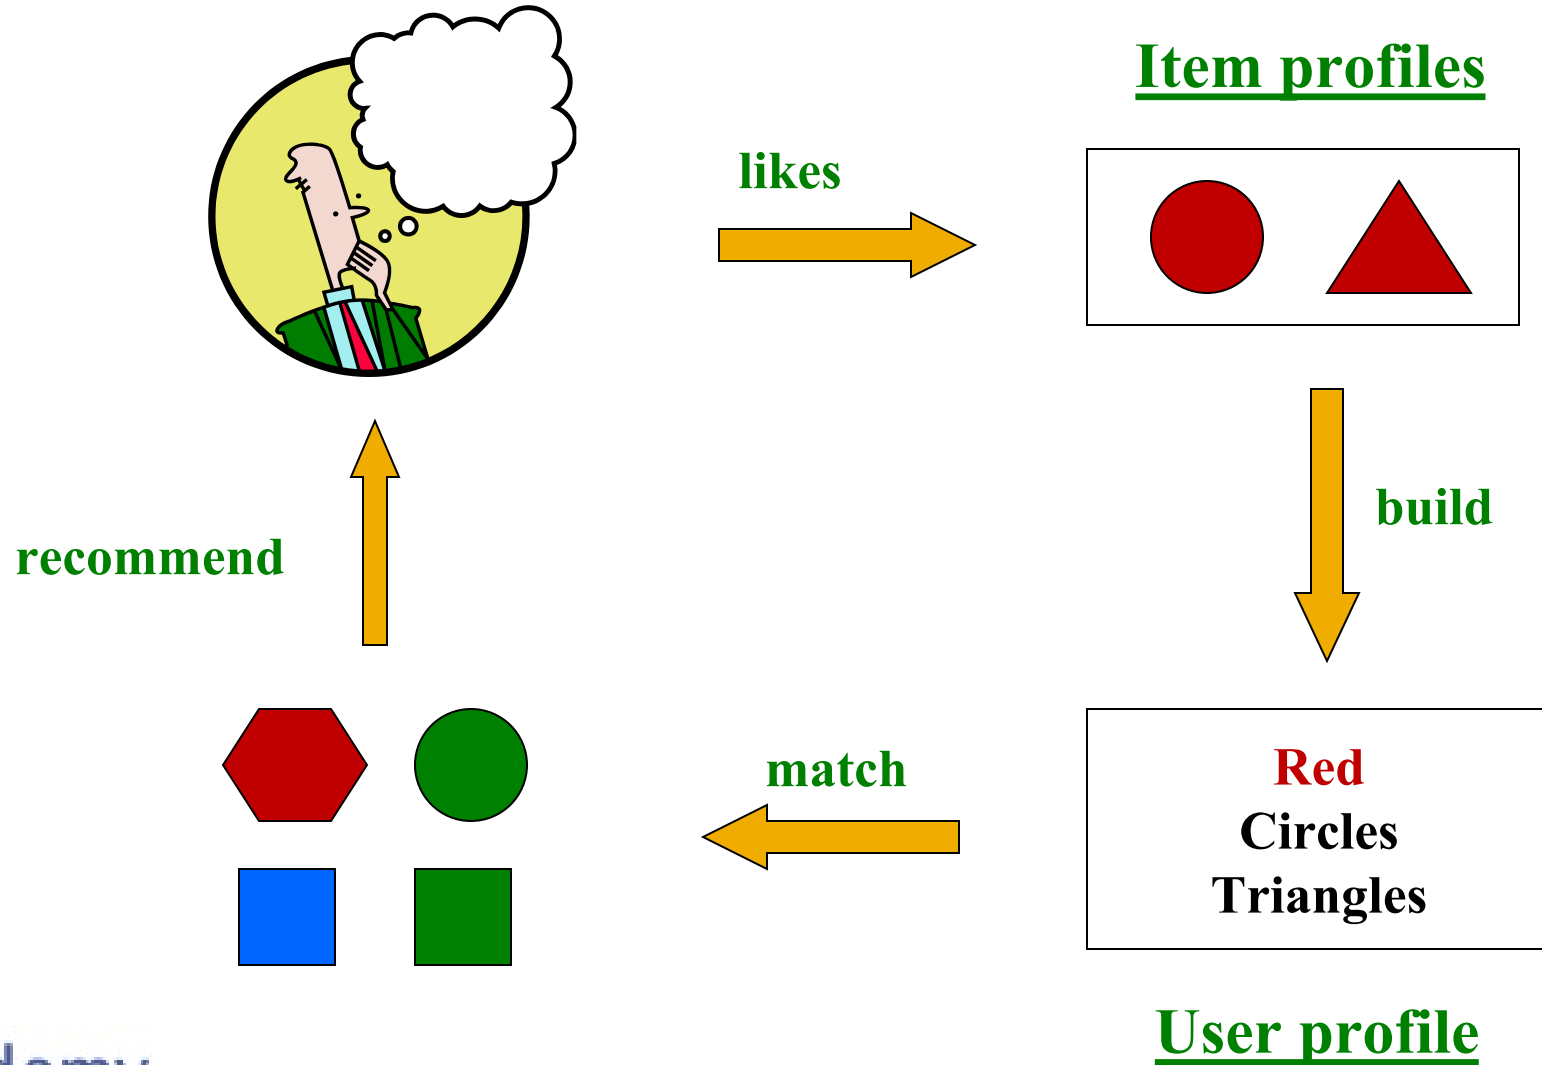
Websites, blogs, video, news

- Articles with “similar” content

People

- People with many common friends.

Plan of Action



Item Profile for Each Item

Profile is a set (vector) of features

- **Movies:** author, title, actor, director,...
- **Text:** Set of “important” words in document
- **People:** Set of friends.

Item profile is a vector

- boolean or real-values

Example: Documents

Profile = set of "important" words in item

How to pick important features?

Usual heuristic from text mining is **TF-IDF**
(Term frequency * Inverse Doc Frequency)

- **Term ... Feature**
- **Document ... Item**

Sidenote: TF-IDF

f_{ij} = frequency of term (feature) i in doc (item) j

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}}$$

n_i = number of docs that mention term i

N = total number of docs

$$IDF_i = \log \frac{N}{n_i}$$

TF-IDF score: $w_{ij} = TF_{ij} \times IDF_i$

Doc profile = set of words with highest **TF-IDF** scores, together with their scores

User Profile

- User has rated items with profiles i_1, \dots, i_n
- (weighted) average of rated item profiles.
- **Variation:** weight by difference from average rating for item
- ...
-

(Weighted) Average of Rated Item Profiles

$[0, 1]$

Items are movies, feature are actor A and actor B

Item Profile: vector with 0 or 1 for each actor

Suppose user X has watched 5 movies:

2 movies featuring actor A

3 movies featuring actor B

$[0.4, 0.6]$

User profile = mean of item profiles

Feature A's weight = $2/5 = 0.4$

Feature B's weight = $3/5 = 0.6$

Weight by Difference from Average Rating for Item²²

Items are movies, feature are actorA and actor B

Item Profile: vector with 0 or 1 for each actor

Suppose user X has watched 5 movies (1-5):

2 movies featuring actor A: rated 3 and 5

3 movies featuring actor B: rated 1, 2 and 4

useful step: **Normalize ratings by subtracting** user's mean rating (3) User profile = mean of item profiles

Feature A's normalized ratings: 0, +2

$$\text{weight} = (0+2)/2 = 1$$

Feature B's normalized ratings: -2, -1, +1

$$\text{weight} = -2/3$$

$$\left[1, \frac{-2}{3} \right]$$

User Profile and Prediction

Prediction heuristic:

- Given user profile \mathbf{x} and item profile \mathbf{i} , estimate

$$u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{\|\mathbf{x}\| \cdot \|\mathbf{i}\|}$$

$$= \frac{1 \cdot 1 - \frac{2}{3} \cdot 0}{\sqrt{1 + \frac{4}{9}} \cdot 1}$$

$$\left[1, -\frac{2}{3} \right]$$

$$\left[1, 0 \right]$$

Pros: Content-based Approach

+: No need for data on other users

+: Able to recommend to users with unique tastes

+: Able to recommend new & unpopular items

- No first-rater problem

+: Able to provide explanations

- Can provide explanations of recommended items by listing content-features that caused an item to be recommended

Cons: Content-based Approach

- : **Finding the appropriate features is hard**
 - E.g., images, movies, music
- : **Recommendations for new users**
 - How to build a user profile?
- : **Overspecialization**
 - Never recommends items outside user's content profile
 - People might have multiple interests
 - Unable to exploit quality judgments of other users

Stay Connected

Dr. Min Chi


Associate Professor

mchi@ncsu.edu

(919) 515-7825



AI Academy



AI Academy

go.ncsu.edu/aiacademy

NC STATE UNIVERSITY